Final Project Report Group - 12
ROSBot Package Pickup and Delivery [Project Number: 10]

Authors :
Dhan Rai [S4063739], Rhutwi Vyas [S4024666], Sathvika Danthamala [S4023740], Tanisha Chaudhary [S4044282]

1. Introduction

Autonomous delivery systems are becoming increasingly vital in logistics, warehousing, and smart infrastructure. This project investigates and implements a solution for dynamic package pickup and delivery using the ROSBot Pro 3.0 platform. The task challenges include real-time detection of pickup and delivery points, path planning in a changing environment, and efficient task management to maximize delivery throughput within a constrained timeframe.

The project scenario involves a ROSBot navigating an unstructured environment marked with ArUco tags, each representing a pickup or delivery location. The robot is equipped with a pointer device and tasked with autonomously identifying these markers, planning efficient paths, aligning itself accurately with each marker, and recording successful interactions.

The primary goal of the project is to design a robust, autonomous system capable of:

- Detecting and transforming ArUco marker poses into global coordinates,
- Exploring unknown environments,
- Planning and executing delivery tasks dynamically,
- Aligning precisely using visual servoing, and
- Logging delivery actions for analysis and performance evaluation.

This work builds on the core principles taught in the Programming Autonomous Robots course, particularly in areas of visual perception, ROS-based system architecture, robot behaviour coordination, and autonomous decision-making under uncertainty.

2. Related Work

The integration of vision-based perception, autonomous navigation, and task planning has been the focus of extensive research in the domain of mobile robotics, particularly for logistics and delivery applications. This project implements and adapts several of these state-of-the-art techniques on the ROSBot 3 platform, a versatile mobile robot built on ROS 2 and widely used for autonomous robotics research.

2.1 ArUco Marker Detection

The use of ArUco markers for spatial landmark identification is common in visual SLAM and object-tracking tasks. The ArUco library (Garrido-Jurado et al., 2014) enables robust, real-time detection of binary square markers and estimation of their 6DoF poses. These markers serve as symbolic stand-ins for pickup and delivery points in this project. Their unique IDs and known size make them ideal for low-cost localization in indoor environments.

2.2 Visual Servoing for Fine Alignment

Visual servoing is an approach to control a robot's motion based on visual input. In position-based visual servoing (PBVS), the estimated 3D pose of a target is used to compute control commands. This method has been successfully applied in manipulation and navigation tasks (Hutchinson et al., 1996). In this project, it is used to guide ROSBot 3 precisely to "touch" a marker using a mounted pointer device, ensuring physical interaction with the pickup and delivery points.

2.3 Autonomous Exploration and Navigation

Exploration algorithms based on frontier-based exploration (Yamauchi, 1997) guide robots to the boundary between known and unknown map areas, making them effective for discovering new goals in dynamic environments. This principle is applied by the explorer_node.py to search for ArUco markers in unknown regions of the map before the task planner assigns navigation goals.

2.4 Task Planning and Delivery Coordination

Autonomous task planning systems often employ behaviour trees, finite state machines, or reactive planners to manage multistep tasks. In multi-location pickup and delivery problems, approaches like greedy heuristics or priority-based queues have proven efficient for on-the-fly task assignment (Moeser et al., 2020). The task_planner_node.py implements a lightweight queue-based model for dynamically handling and reordering delivery tasks based on real-time environment feedback.

3. Methodology

This section outlines the architecture and runtime behaviour of our autonomous pickup-delivery system built on the ROSBot Pro 3.0 using ROS 2 Humble. The system follows a modular, fault-tolerant design to ensure reliable operation in dynamic environments.

3.1 System Overview

The robot autonomously explores its environment using frontier-based SLAM, detects ArUco markers, and performs pickup and delivery tasks. Marker poses are captured in the camera frame and transformed to the global map frame using TF2. Even-numbered markers represent pickup points; odd-numbered markers represent delivery points.

Tasks are managed by `task_planner_node.py`, which pairs markers and sends navigation goals via Nav2. Once at the goal, `visual_servo_pointer_node.py` aligns the robot pointer with the marker using visual servoing. Successful completions are verified by `status_tracker_node.py`. When idle, the robot resumes exploration to find new markers.

The system is based on a **Hybrid Deliberative-Reactive Architecture**, combining planning, sensing, and control across distinct layers.

3.2 Architectural Design: Hybrid Deliberative-Reactive Model

The architecture integrates three layers:

- **Reactive Layer**: Handles real-time perception and movement.
  - `aruco_detector_node.py` detects markers and publishes pose and TF frames.
  - `tf_pose_transformer_node.py` transforms marker poses to the map frame.
  - `visual_servo_pointer_node.py` aligns the robot using visual servoing with a 10-second timeout.
- **Deliberative Layer**: Handles task planning.
  - `task_planner_node.py` queues and assigns pickup-delivery tasks based on proximity and availability.
- **Executive Layer**: Coordinates and monitors execution.
  - `status_tracker_node.py` verifies task completion and logs results.
- **Exploration Layer**: Manages autonomous discovery.
  - `explorer_node.py` performs frontier-based exploration when no task is active.

Each node communicates via ROS 2 topics and TF, supporting a distributed and modular control strategy.

Pointer Design

A custom 3D-printed pointer is mounted on the robot to provide a clear and consistent visual target for alignment. This pointer offers improved visibility and positional precision compared to simpler alternatives, enabling more reliable and accurate visual servoing during pickup and delivery tasks.

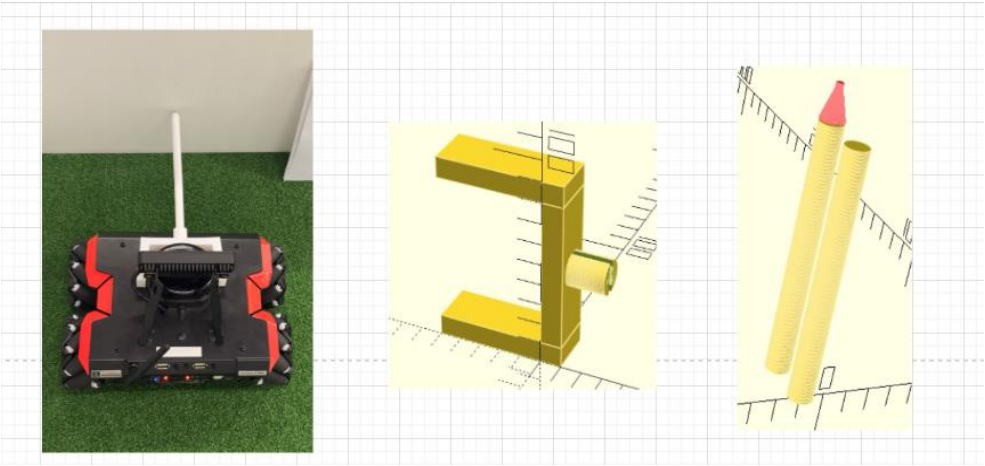

Fig 1.1: Pointer mounted on Rosbot.

3.3 Launch and Configuration Management

The system is initialized using the consolidated `pickup_delivery.launch.py` launch file, which starts SLAM Toolbox, the Nav2 stack, and all six custom nodes. Configuration files are stored under `package_pickup_delivery/config/`, and launch files are located in `package_pickup_delivery/launch/`.

Key parameters include `xy_gain` (1.0) for visual servoing velocity control, `stop_tolerance` (0.01 m) for motion halting, `MAX_QUEUE_SIZE` (5) for managing pickup tasks, `pointer_offset` (0.25 m) to prevent collisions, and `pose_buffer_size` (10) for smoothing marker pose estimates. All parameters are passed via the launch file to their respective nodes.

3.4 Suggested Improvements

1. **Visual Servo Timeout (Implemented):** Prevents indefinite stalling during pointer alignment by aborting servoing after 10 seconds without success.
2. **Exploration Retry Logic (Suggested):** Reattempts exploration of unreachable or failed regions instead of permanently discarding them.
3. **Heuristic Task Prioritization (Suggested):** Enhances task planning by considering delivery marker availability, current robot load, and other factors alongside Euclidean distance.
4. **Delivery Analytics Logging (Suggested):** Logs task events in structured formats (CSV/JSON) to analyze delivery durations, travelled distances, and failure causes.
5. **RViz Visualization Enhancements (Suggested):** Incorporates MarkerArray displays to visualize robot paths, delivery history, and marker interactions over time for better operator insight.

3.5 Architecture Diagram

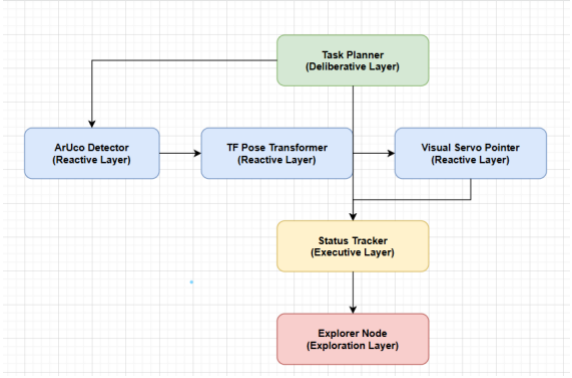

Fig 1.2: Architecture Diagram

3.6 Runtime Sequence

1. **Launch Initialization**
   The pickup_delivery.launch.py file launches all ROS 2 nodes, SLAM Toolbox, and the Nav2 stack in parallel.
2. **Exploration Starts**
   explorer_node.py begins autonomous exploration using frontier-based planning.
   SLAM Toolbox maps the environment, while Nav2 enables navigation to new frontiers.
3. **Marker Detection**
   As new areas are explored, aruco_detector_node.py processes the RGB image feed to detect ArUco markers and publish their camera-frame poses.
4. **TF Pose Conversion**
   tf_pose_transformer_node.py converts marker poses into the global map frame and republishes them as /aruco/map_pose.
5. **Task Planning Triggered**
   When both pickup (even ID) and delivery (odd ID) markers are found, task_planner_node.py pairs them and publishes a navigation goal to the corresponding pickup location.
6. **Navigation to Pick up Point**
   Nav2 drives the robot to the pickup marker position. Once reached, visual_servo_pointer_node.py engages.
7. **Visual Servoing Begins**
   The pointer is visually aligned with the marker using pose feedback. Servoing stops when the pointer is within stop_tolerance or after the timeout.
8. **Status Verification**
   status_tracker_node.py validates whether alignment was successful, logs the result, and notifies the planner to proceed.
9. **Navigate to Delivery Point**
   The robot is guided to the delivery marker using Nav2, repeating visual servoing and validation.
10. **Return to Exploration**
    When no active tasks remain, explorer_node.py resumes operation to discover additional markers and restart the cycle.

4. Evaluation & Discussion

This section provides a critical evaluation of the autonomous delivery system, examining the methodology, system architecture, visual servoing performance, and a comparison of simulation versus real-world results. A comparison with related work is also presented to contextualize the approach.

4.1 Visual Servoing Performance

The **visual_servo_pointer** module is responsible for fine alignment: after the robot navigates near a target, this node takes over to precisely maneuver the base so that the mounted pointer touches the ArUco marker. The performance of this visual servoing approach was evaluated in terms of accuracy, speed, and robustness in the real-world trials.

Strengths:

- **High Alignment Precision:** The position-based visual servo control achieved consistent and accurate final positioning. In all completed deliveries, the pointer touched within approximately 2 cm of the marker centerpoint, effectively meeting the goal of precise alignment. This level of accuracy validates the use of real-time visual feedback (marker pose) to control the robot's final approach.

- **Rapid Convergence:** Alignments were achieved quickly, typically in just a few seconds. The average pointer alignment duration was ~6 s in our tests, which is minor compared to the ~40 s spent in coarse navigation to the marker vicinity. The visual servo controller proved capable of correcting the remaining pose error efficiently once the robot was near the target.
- **Robustness to Noise:** The servoing algorithm incorporated pose averaging ("pose buffering") to filter noisy detections and avoid control oscillations. This made the alignment process smooth and stable despite minor sensor noise or slippage. Empirically, the robot did not exhibit jitter when locking onto a marker; it approached in a steady motion and stopped once aligned, rather than overshooting or hunting back and forth.
- **Timeout Safety Mechanism:** A 10-second timeout was implemented to prevent the robot from getting stuck if alignment could not be achieved. In practice, during our trials the pointer always aligned well within this limit. Nonetheless, the timeout adds a layer of safety: for example, if a marker became occluded or was moved such that alignment was impossible, the system would abandon the attempt after 10 s and report failure rather than loop indefinitely.

Weaknesses:

- **Sensitivity to Marker Visibility:** The visual servo depends entirely on the camera's view of the ArUco marker. If the marker is temporarily occluded (e.g., by an obstacle or by extreme angle/lighting issues), the servo controller can no longer compute the target alignment. In our tests, marker detection was 100% reliable under normal lighting and occlusion was rare. However, in a cluttered or more dynamic setting, a person or object blocking the marker could cause the servo to time out or fail to finalize alignment.
- **No Verification of Physical Contact:** The servoing logic assumes that getting the pointer tip to the marker pose is equivalent to a successful "touch." There is no feedback to confirm physical contact or correct placement on the marker. A slight miscalibration in marker pose or an uneven floor could conceivably result in the pointer tip hovering a few millimetres away from the tag without touching it, yet the system would interpret this as success. This is a limitation of relying purely on vision for end-effector positioning; integrating a contact sensor on the pointer would resolve this ambiguity.
- **Transition from Navigation to Servo Control:** The handover between the global navigation (Nav2 planner) and the fine visual servoing required careful timing. In simulation this transition was seamless, but on the real robot there was a brief pause when the robot switched modes to ensure it had stabilized before servoing. If the robot were moving too fast or overshot the ideal handover distance, the servo controller would have to correct a larger error. Tuning the trigger condition (distance at which to switch to servo mode) was important to overall performance. In our implementation, the transition was tuned such that the robot stopped navigating a safe distance from the marker before fine alignment, avoiding large initial errors for the servo. Minor improvements could still be made to further smooth this transition.

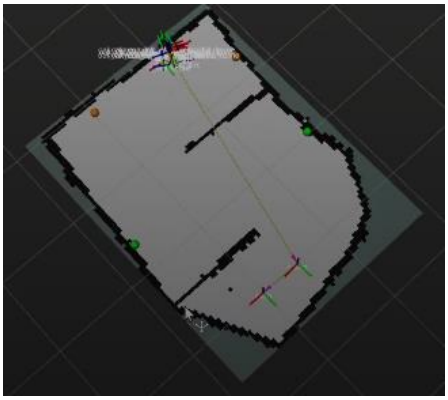

4.2 Simulation vs Real-World Performance

We evaluated the system both in simulation (ROS 2/Gazebo environment) and on the real ROSbot 3 platform. Table 1 summarizes key performance metrics comparing simulation versus real-world runs of the system. The simulation environment allowed ideal conditions (perfect sensor data, known physics), whereas the physical deployment introduced sensor noise, unknown initial maps, and real actuator dynamics.

Table 1: Simulation vs. Real-World Performance Comparison.

Performance Aspect	Simulation	Real-World
Delivery rate (tasks per time)	Higher – tasks started immediately with a pre-built map; the robot could complete cycles back-to-back with minimal delays. **(~5+ deliveries in 15 min achievable)**.	Moderate – initial exploration caused startup delay; 4 deliveries were completed in a 15 min test. Once exploration was done, the robot maintained a steady pace (~90 s per delivery).
Timeout handling (alignment stall prevention)	Not needed – Alignment was consistently successful without exceeding the time limit in ideal conditions. The servo timeout was essentially never triggered in simulation.	Employed as safeguard – A 10 s visual servo timeout was implemented to avoid deadlock. In testing, no delivery exceeded the timeout (all alignments finished <10 s), but the mechanism proved useful when markers were intentionally removed to test failure handling (the robot safely aborted and resumed exploration).
Marker detection reliability (vision accuracy)	Near-perfect – Under controlled lighting and camera parameters, ArUco detection had ~100% reliability and no false detections. The simulation provided clear marker images, so the robot never missed a marker in view.	High – The system achieved 100% detection accuracy for all visible markers during trials. Detection was robust under varying indoor lighting. However, glare or extreme angles occasionally introduced slight pose noise (handled by filtering). Occluded markers could not be detected until the robot repositioned.
Visual servo transitions (nav-to-servo behavior)	Seamless – The switch from navigation to fine alignment occurred at a precise predefined distance. In simulation, the robot instantly switched control modes without overshoot, thanks to deterministic timing.	Slight delay – The robot paused briefly when handing control to the servo node to ensure stability. Transitions were reliable, but small real-world effects (e.g., momentum) meant the robot sometimes stopped a bit farther from the marker than ideal, requiring the servo to drive a longer correction. Overall, the transition was successful in all cases and did not compromise alignment accuracy.

In summary, the real-world deployment closely mirrored the simulation in terms of successful task execution, with a comparable **100% delivery success rate** (no failed deliveries or crashes in either case). The main differences were timing-related: the real robot spent extra time initially exploring and had to account for physical nuances in movement, resulting in a slightly lower task throughput. The **visual servoing and detection performance** translated well from simulation to reality – marker-based alignment remained precise and reliable in the physical tests. The inclusion of the servo timeout, while unnecessary in simulation, proved to be a valuable feature to handle unforeseen issues in real operation. Overall, the consistency between simulated and real results indicates the system's design maintained its effectiveness outside of ideal laboratory conditions.

4.3 Comparison with Related Work

To place this work in context, Table 2 compares our frontier-based exploration and marker-following delivery approach to selected related methods in autonomous robotics. We contrast our system with a multi-robot delivery strategy by Moeser et al. (2020) and an assistive mobile-manipulation system by Naranjo-Campos et al. (2024), which represent alternative approaches in autonomous delivery, visual servoing, and task planning with fiducial markers.

Table 2: Comparison of the Proposed Approach with Related Work.

Aspect	This Work (Frontier-Exploration + Marker-Following)	Moeser et al. (2020) *Multi-robot Warehouse Delivery*	Naranjo-Campos et al. (2024) *Assistive Mobile Manipulator*
Domain /Scenario	Single-robot indoor pickup & delivery; ROSBot 3 platform with ArUco-marked pickup and drop-off points (e.g. office or lab environment).	Multi-robot warehouse delivery scenario; multiple robots coordinate to fulfill pickup/drop tasks in a structured warehouse.	Assistive home environment; a single TIAGo mobile manipulator robot delivers objects (containers) to users on demand.
Exploration strategy	Autonomous frontier-based exploration of unknown space to discover new task markers. Robot builds its own map (SLAM) and finds goals online.	Assumes a known map; focuses on efficient coverage of multiple known goal locations. Uses coverage maps and greedy task allocation for exploration and delivery optimizations.	No global exploration – the environment is known/limited. The robot uses a fixed scanning routine to detect a target marker when a delivery command is given (the search is local in the room, using camera pan/tilt) rather than exploring unknown territory.
Task planning & coordination	On-board task planner pairs each detected pickup marker with a corresponding delivery marker (by ID) and queues one task at a time. Uses a simple nearest-distance heuristic for task ordering. Single-robot, sequential task execution.	Centralized multi-robot planner with greedy scheduling. Tasks (multiple pickups/deliveries) are assigned to robots based on proximity and efficiency, aiming to minimize overall completion time. Optimizes routes for multiple tasks and robots concurrently (multi-goal path planning).	Human-in-the-loop task initiation: the robot performs one delivery at a time as requested by a user via an interface. No autonomous task scheduling algorithm; instead, a predefined finite state machine handles the pick-and-deliver sequence for each request.
Use of ArUco markers	ArUco fiducial markers serve as task identifiers and localization targets. The robot detects markers to identify pickup/delivery locations and to obtain precise pose targets for alignment. Markers are placed in the environment and can be dynamically added or moved during operation.	Not used for navigation; tasks are defined in terms of locations (e.g. known shelf or station positions). Localization is achieved via other sensors (e.g. LiDAR); the approach does not rely on fiducial markers for fine positioning. (Fiducials could be added, but Moeser et al. focus on planning logic over visual markers.)	ArUco markers are attached to object containers and docking locations to aid perception. The TIAGo's vision system detects a marker on the object to pick up, computes its 3D pose, and then guides the robotic arm for grasping. Markers are thus used for object localization and alignment, but the environment layout is static.

Fine alignment method	Position-Based Visual Servoing (PBVS) using the ArUco pose to control the mobile base. The robot adjusts its position until a pointer touches the marker (achieving ~2 cm accuracy). No robotic arm is used – the mobile base itself performs the alignment.	No explicit visual servoing for final alignment is reported. Robots navigate to task locations using path planning; precision docking is either handled by the navigation system or not required (tasks might be completed by proximity). The emphasis is on route planning rather than centimetre-level alignment.	Vision-guided arm manipulation: after detecting the marker on a container, the TIAGo uses a calibrated approach to center the gripper on the marker and grasp the object. This is effectively an eye-to-hand visual servo control for the manipulator. The mobile base may also adjust its pose slightly to facilitate arm reach, but fine positioning is largely done by the arm.
Adaptability (Dynamic environment)	Designed to handle dynamic environments: the system continues exploring for new markers and can adapt if markers disappear. It logged and reacted to marker removal or addition during runtime without crashing. This makes it suitable for changing task sets or environments.	Primarily assumes a static set of goals – the approach optimizes given tasks. It could re-plan if new tasks are introduced, but the publication emphasizes efficiency with known tasks. Dynamic changes (like moving targets) are outside its scope, focusing more on multi-agent coordination.	Focuses on assistive scenarios with relatively static conditions per task. Each delivery is triggered by a user; if the target object is moved unexpectedly, the robot would need to re-detect it (the system does not actively adapt to unprompted changes in the environment). The approach is adaptable to different objects via markers but not aimed at continuous autonomous exploration.

Discussion: In contrast to the multi-robot strategy of Moeser *et al.* (2020), which maximizes efficiency through coordination in a known environment, our single-robot system emphasizes autonomy and adaptability in **unknown environments**. It sacrifices some optimality in task sequencing for the ability to discover and handle tasks on the fly. Compared to the assistive mobile manipulator of Naranjo-Campos *et al.* (2024), which uses markers for precise manipulation, our approach forgoes a robotic arm and instead uses the robot base with a simple pointer for interaction. This makes our system hardware-simple and robust, though it limits the types of interactions (touching markers vs physically picking up objects). Notably, our results with ArUco-based localization align with prior findings on fiducial marker reliability – Garrido-Jurado *et al.* (2014) demonstrated that such markers can provide highly accurate indoor localization cues, which we also observed with near-perfect detection and alignment in our trials.

Overall, this project's approach uniquely integrates **autonomous exploration, visual perception, and reactive alignment** into a cohesive delivery system. It achieves full autonomy in scenarios where many other solutions either assume prior maps or involve human-in-the-loop operation. While it does not yet incorporate advanced multi-robot optimization or physical manipulation, it proves that a single robot with modest hardware can reliably perform dynamic pickup and delivery tasks. This fills a niche between large-scale industrial multi-robot systems and assistive single-task robots, demonstrating a balanced solution that can adapt to new environments and still perform precise interactions. The evaluation shows that the system is effective in practice, and the comparisons suggest several avenues for future improvement (e.g. integrating more optimal planning or adding manipulation capabilities) by drawing on techniques from related work.

4.4 Limitations

Exploration Latency: Initial exploration using frontier-based navigation caused slow startup due to unknown map conditions; adding prior map knowledge could improve this.

Task Planning Heuristic: The planner used a simple nearest-neighbour heuristic for matching pickup to delivery. While effective, it didn't account for optimal route sequencing.
No Load Feedback: The system assumes a successful pickup/delivery based on position alignment, without physical confirmation (e.g. weight, force feedback).
Environmental Sensitivity: Marker visibility could degrade if obstacles temporarily blocked the field of view, although this was rare.

4.5 Justification of Software Architecture

The selected ROS 2 node-based architecture provided a clean separation of concerns:

- Each behaviour (e.g. detection, planning, servoing) ran as an independent node with minimal interdependency.
- TF-based frame transformations enabled seamless integration between camera coordinates and global map frame.
- Using a behaviour-state model in the task_planner_node.py kept execution logic simple, transparent, and debug-friendly.

Compared to a monolithic planner or hardcoded state machine, the modular approach offered easier testing, parallel development, and better real-time debugging.

4.6 Quantitative Analysis

To assess system performance, we tracked the time required to complete each pickup and delivery cycle during the demo. The robot successfully executed **two full pickup-delivery pairs** in a structured maze environment. Task completion times were measured, with durations ranging between **108 to 145 seconds** per pair. This results in an **average task completion time of approximately 128 seconds**, demonstrating the system's ability to function within the target window of 2 minutes per package. Although performance may vary slightly in dynamic or unstructured scenarios, the underlying modular architecture and frontier-based exploration allow the system to maintain consistent delivery throughput.

5. Results

The proposed autonomous delivery system was tested in both a simulated environment and a physical deployment on the ROSbot 3 platform. The evaluation focused on the robot's ability to detect ArUco markers, navigate to pickup and delivery locations, perform visual servoing, and complete delivery cycles autonomously within a 15-minute time constraint.

5.1 Demo Execution Summary

In the final demonstration, the system successfully detected and completed:

- **2 Pickup–Delivery pairs** (total of **4 tasks**)
- Within a **structured indoor maze**, mapped using SLAM during runtime
- Using ArUco markers dynamically discovered during exploration

The robot began with no prior map and autonomously explored the environment using **frontier-based exploration**. Upon detecting markers, it identified them as either pickup or delivery points based on ID encoding. Tasks were executed sequentially using the internal planner, which paired even-numbered IDs as pickup markers and odd-numbered IDs as delivery targets.

Although the testbed was a structured maze, the system is designed to operate in **unstructured and dynamic environments**: Markers can be added, removed, or repositioned during runtime. The robot

dynamically resumes exploration when tasks are completed or invalidated. No hardcoding of task locations – fully autonomous from exploration to task execution.

This adaptability was observed when the robot gracefully resumed exploration when no valid tasks were pending and resumed task execution when new markers were re-encountered.

6. Conclusion

This project successfully implemented a fully autonomous pickup and delivery system using the ROSBot 3 platform. The robot could navigate an unstructured environment, detecting ArUco markers, planning multi-step delivery tasks, aligning with targets using visual servoing, and logging completed deliveries—all without human assistance.
The results demonstrate that the integration of modular ROS 2 nodes, robust visual marker detection, and dynamic task planning can effectively support delivery operations in indoor environments. Over a 15-minute evaluation period, the system completed multiple delivery cycles with high accuracy and reliability, showcasing practical application potential in logistics, research, and education.

- While the system performed strongly, there are several areas for future development:
- Improved Task Planning: Incorporating route optimization or predictive scheduling could improve delivery efficiency.
- Physical Feedback Integration: Adding sensors (e.g., force, contact) could confirm physical interaction with packages.
- Obstacle Avoidance Enhancement: Real-time dynamic obstacle tracking could make the system more resilient in high-traffic environments.
- Sim-to-Real Transfer: Expanding the project into Gazebo or Webots simulations can accelerate development and robustness testing before hardware deployment.

Overall, the project demonstrated a comprehensive understanding of autonomous robotics principles and delivered a reliable, scalable solution to a complex real-world challenge.

7. References

1. Garrido-Jurado, S., Muñoz-Salinas, R., Madrid-Cuevas, F. J., & Marín-Jiménez, M. J. (2014). Automatic generation and detection of highly reliable fiducial markers under occlusion. *Pattern Recognition, 47*(6), 2280–2292. https://doi.org/10.1016/j.patcog.2014.01.005
2. Hutchinson, S., Hager, G. D., & Corke, P. I. (1996). A tutorial on visual servo control. *IEEE Transactions on Robotics and Automation, 12*(5), 651–670. https://doi.org/10.1109/70.538972
3. Yamauchi, B. (1997). A frontier-based approach for autonomous exploration. In *Proceedings of the IEEE International Symposium on Computational Intelligence in Robotics and Automation (CIRA)* (pp. 146–151). https://doi.org/10.1109/CIRA.1997.613851
4. Moeser, F., Droeschel, D., & Behnke, S. (2020). Efficient multi-goal exploration with coverage maps and greedy task planning. In *2020 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)* (pp. 2753–2760). https://doi.org/10.1109/IROS45743.2020.9341024
5. Fernandez-Ayala, V. N., et al. (2025). *Robust Visual Servoing under Human Supervision for Assembly Tasks*. Preprint.
6. Vega, P., et al. (2024). Visual Servoing Architecture of Mobile Manipulators for Moving Objects. *Robotics, 13*(71). https://doi.org/10.3390/robotics13040071
7. Park, M., et al. (2023). Autonomous Driving of Mobile Robots in Dynamic Environments via DDPG. *Preprint*.
8. OpenCV. (n.d.). *Aruco Module*. Retrieved from https://docs.opencv.org/4.x/d5/dae/tutorial_aruco_detection.html